

## Data Preparation- Pandas

### Series

A Series is very similar to a NumPy array . What differentiates the NumPy array from a Series, is that a Series can have axis labels, meaning it can be indexed by a label, instead of just a number location. It also doesn't need to hold numeric data, it can hold any arbitrary Python Object.

```
In [124]: 1 import numpy as np
          2 import pandas as pd
          3 #from pandas import Series, DataFrame
```

```
In [125]: 1 Series_obj = pd.Series(np.arange(8), index=['row 1', 'row 2', 'row 3', 'row 4', 'row 5', 'row 6', 'row 7', 'row 8'])
          2 Series_obj
```

```
Out[125]: row 1    0
          row 2    1
          row 3    2
          row 4    3
          row 5    4
          row 6    5
          row 7    6
          row 8    7
          dtype: int32
```

Now we want to select an element with the label index of row 7:

```
In [67]: 1 Series_obj['row 7']
```

```
Out[67]: 6
```

### DataFrame

Create a DataFrame object:

Here is an example of 36 random number in a 6x6 matrices.

```
In [126]: 1 np.random.seed(25)
          2 DF_obj = pd.DataFrame(np.random.rand(36).reshape((6,6)),
          3                        index=['row 1', 'row 2', 'row 3', 'row 4', 'row 5', 'row 6'],
          4                        columns=['column 1', 'column 2', 'column 3', 'column 4', 'column 5', 'column 6'])
          5 DF_obj
```

```
Out[126]:
```

	column 1	column 2	column 3	column 4	column 5	column 6
row 1	0.870124	0.582277	0.278839	0.185911	0.411100	0.117376
row 2	0.684969	0.437611	0.556229	0.367080	0.402366	0.113041
row 3	0.447031	0.585445	0.161985	0.520719	0.326051	0.699186
row 4	0.366395	0.836375	0.481343	0.516502	0.383048	0.997541
row 5	0.514244	0.559053	0.034450	0.719930	0.421004	0.436935
row 6	0.281701	0.900274	0.669612	0.456069	0.289804	0.525819

### Reading csv file

```
In [ ]: 1 #pd.read_csv
```

```
In [3]: 1 import pandas as pd
2
3 #How we read in a pandas dataframe. The header=0 means column names are in the first row
4 df=pd.read_csv('Downloads/mtcars.csv', header=0)
5 df
6
```

Out[3]:

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
5	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
6	Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
7	Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
8	Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
9	Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
10	Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
11	Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
12	Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
13	Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
14	Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
15	Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
16	Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
17	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
18	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
19	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
20	Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
21	Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
22	AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
23	Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
24	Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
25	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
26	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

## The head() Method

```
In [4]: 1 #The head method returns the first five rows
2 df.head()
```

Out[4]:

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

## Basic Features

```
In [76]: 1 #There are column names
        2 df.columns
```

```
Out[76]: Index(['name', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am',
               'gear', 'carb'],
              dtype='object')
```

```
In [77]: 1 #And there are row names
        2 list(df.index)
```

```
Out[77]: [0,
          1,
          2,
          3,
          4,
          5,
          6,
          7,
          8,
          9,
          10,
          11,
          12,
          13,
          14,
          15,
          16,
          17,
          18,
          19,
          20,
          21,
          22,
          23,
          24,
          25,
          26,
          27,
          28,
          29,
          30,
          31]
```

```
In [79]: 1 #Get the dimensions of the data frame with shape
        2 dimensions = df.shape
        3 dimensions
```

```
Out[79]: (32, 12)
```

```
In [80]: 1 #Get the data type of each column
        2 df.dtypes
```

```
Out[80]: name      object
        mpg    float64
        cyl      int64
        disp    float64
        hp       int64
        drat    float64
        wt       float64
        qsec    float64
        vs       int64
        am       int64
        gear    int64
        carb    int64
        dtype: object
```

```
In [81]: 1 #We can pick out a column by referencing its name. The result is a series or one dimensional data frame
        2 df['mpg'].head()
```

```
Out[81]: 0    21.0
          1    21.0
          2    22.8
          3    21.4
          4    18.7
        Name: mpg, dtype: float64
```

```
In [82]: 1 #You can similarly pick out columns as attributes with the '.'
         2 df.mpg.head()
```

```
Out[82]: 0    21.0
         1    21.0
         2    22.8
         3    21.4
         4    18.7
         Name: mpg, dtype: float64
```

Note that when we slice a series, the second entry is non-inclusive.

```
In [83]: 1 #You can pick out multiple columns by specifying a list of column names
         2 name_grade = df[['mpg', 'cyl', 'hp']].head()
         3 name_grade
```

```
Out[83]:      mpg  cyl  hp
0    21.0   6  110
1    21.0   6  110
2    22.8   4   93
3    21.4   6  110
4    18.7   8  175
```

## Slicing and Indexing

We will be using the .loc (just labels) approach. You can also slice with .iloc (just indices) or .ix (indices and labels).

```
In [84]: 1 #Let's Look at the data
         2 df.head()
```

```
Out[84]:      name  mpg  cyl  disp  hp  drat   wt  qsec  vs  am  gear  carb
0  Mazda RX4   21.0   6  160.0  110  3.90  2.620  16.46  0   1    4    4
1  Mazda RX4 Wag 21.0   6  160.0  110  3.90  2.875  17.02  0   1    4    4
2    Datsun 710  22.8   4  108.0   93  3.85  2.320  18.61  1   1    4    1
3  Hornet 4 Drive 21.4   6  258.0  110  3.08  3.215  19.44  1   0    3    1
4  Hornet Sportabout 18.7   8  360.0  175  3.15  3.440  17.02  0   0    3    2
```

```
In [85]: 1 #Pick out a single entry
         2 df.loc[3, "name"]
```

```
Out[85]: 'Hornet 4 Drive'
```

```
In [87]: 1 #Select contiguous rows and columns
         2 df.loc[1:5, "disp": "wt"]
```

```
Out[87]:      disp  hp  drat   wt
1  160.0  110  3.90  2.875
2  108.0   93  3.85  2.320
3  258.0  110  3.08  3.215
4  360.0  175  3.15  3.440
5  225.0  105  2.76  3.460
```

```
In [88]: 1 #Select none contiguous rows
         2 df.loc[[0,2,4], ["wt", "am"]]
```

```
Out[88]:      wt  am
0  2.62   1
2  2.32   1
4  3.44   0
```

## Built in Functions

- Useful built in column methods.
- Creating new columns and deleting existing ones.

```
In [89]: 1 #Read in the data frame
2 df=pd.read_csv("Downloads/mtcars.csv", header=0)
3
4 df.head()
```

```
Out[89]:
```

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

```
In [90]: 1 #Compute mean of carb column
2 avg_final = df["carb"].mean()
3 avg_final
```

```
Out[90]: 2.8125
```

## Creating New Columns

Next, we look at how to create new columns

```
In [91]: 1 #Create a New Column that is a function of other columns
2 df["carb_new"] = df["carb"]/2
3 df.head()
```

```
Out[91]:
```

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	carb_new
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	2.0
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4	2.0
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1	0.5
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1	0.5
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	1.0

## Deleting Columns (Drop Method)

```
In [92]: 1 #I can then delete it with the drop method
2 df.drop(["carb_new"], inplace = True, axis=1)
3 df.head()
```

```
Out[92]:
```

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

The inplace argument works as follows:

- inplace = True : The dataframe itself will have the given column(s) deleted.
- inplace = False: Will return a dataframe with the column(s) deleted.

The axis argument works as follows:

- axis = 1 : delete columns given
- axis = 0 : delete rows given.

Let's look at an example where we delete rows

```
In [93]: 1 #Delete rows with index 0 and 2
         2 drop_rows = df.drop([0,2], inplace = False, axis=0)
         3 drop_rows.head()
```

Out[93]:

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
5	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
6	Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4

Let's have a look at df

```
In [94]: 1 df.head()
```

Out[94]:

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

Note that df was not changed! This is what happens when you set inplace.

Let's see how we can sort a data frame. The inplace argument has the same affect as the drop method.

```
In [95]: 1 #Sort the data frame according tothe mpg Column
         2 #By setting inplace= False will just return the sorted dataframe and not chnage df
         3 df.sort_values(by = ["mpg"], inplace =False, ascending=False).head()
```

Out[95]:

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
19	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
17	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
18	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
25	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1

Now let's sort by multiple columns, specifying more than one column is essentially specifying a tie break

```
In [97]: 1 #Sort by Mini Exam 1 and tie breal with Previous Part
         2
         3 result_sorted = df.sort_values(by = ["mpg", "wt"], inplace =False, ascending=False)
         4 result_sorted.head()
```

Out[97]:

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
19	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
17	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
18	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
25	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1

In this part, we will a collection of important miscellaneous concepts that include:

- Changing columns names
- Combining dataframes
- Understanding the index
- Missing Data

```
In [99]: 1 import pandas as pd
2
3 #Read in the data frame
4 df=pd.read_csv("Downloads/mtcars.csv", header=0)
5
6 df.head()
```

```
Out[99]:
```

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

Recall that we can get the column names through the attribute column

```
In [100]: 1 #Get the column names
2 df.columns
```

```
Out[100]: Index(['name', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am',
'gear', 'carb'],
dtype='object')
```

## Changing Column Names

We can change column names through the rename method

```
In [101]: 1 #Change the column names
2 df.rename(columns={"mpg": "mpg_1", "cyl": "cyl_1"}, inplace=True)
3
4 df.head()
```

```
Out[101]:
```

	name	mpg_1	cyl_1	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

## Concatenation

```
In [ ]: 1 #pd.concat
```

Next, we see how to combine or concatenate two (or more) data frames.

```
In [102]: 1 #I can combine data frames with concat function
2 head = df.head()
3 tail = df.tail()
4
5
```

```
In [103]: 1 #Have a look at the variable head
2 head
```

```
Out[103]:
```

	name	mpg_1	cyl_1	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2

```
In [104]: 1 #Have a look at the variable head
          2 tail
```

```
Out[104]:
```

	name	mpg_1	cyl_1	disp	hp	drat	wt	qsec	vs	am	gear	carb
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.9	1	1	5	2
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.5	0	1	5	4
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.5	0	1	5	6
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.6	0	1	5	8
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.6	1	1	4	2

```
In [105]: 1 #axis=0 says stack them top to bottom. axis =1 stacks side to side
          2 dfConcat = pd.concat([head,tail], axis =0)
          3 dfConcat
```

```
Out[105]:
```

	name	mpg_1	cyl_1	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

## Handling Missing Data

Missing data is common in most data analysis applications. You have a number of options for filtering out missing data. One option is doing it by hand or you can use the *dropna* method.

With dataframes objects, things get a little more complex. You may want to drop rows or columns which are all NA or just those containing any NAs. *dropna* by default drops any row containing a missing value.

```
In [106]: 1 #Here we have two pieces of missing data
          2 df_missing = pd.read_csv("Downloads/mtcars_missing.csv")
          3 df_missing
```

```
Out[106]:
```

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160	110.0	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	NaN	6	160	110.0	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108	93.0	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258	NaN	3.08	3.215	19.44	1	0	3	1

The `isnull()` method returns a series or dataframe of booleans corresponding to whether the particular entries are null or not.

```
In [107]: 1 #isnull method for a data frame
          2 df_missing.isnull()
```

```
Out[107]:
```

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	False	False	False	False	False	False	False	False	False	False	False	False
1	False	True	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	True	False	False	False	False	False	False	False

We can make sure they are all read in as NA values using the `na_values` input when we read in the file

Now lets see how we can change/replace these NA values



```
In [109]: 1 #Get rid of all rows with an NA
          2 df_missing.dropna(axis=0)
```

```
Out[109]:
```

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160	110.0	3.90	2.62	16.46	0	1	4	4
2	Datsun 710	22.8	4	108	93.0	3.85	2.32	18.61	1	1	4	1

Rather than filtering out missing data, you may want to fill in the "holes" in any number of ways. For most purposes, the *fillna* method with a constant replaces missing values with that value.

```
In [110]: 1 df_missing.fillna(0)
```

```
Out[110]:
```

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160	110.0	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	0.0	6	160	110.0	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108	93.0	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258	0.0	3.08	3.215	19.44	1	0	3	1

```
In [112]: 1 #You can pass fillna a dict which gives the replacement value for each column
          2 df_missing.fillna({"mpg":20,"hp":100})
```

```
Out[112]:
```

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160	110.0	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	20.0	6	160	110.0	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108	93.0	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	258	100.0	3.08	3.215	19.44	1	0	3	1

With *fillna* you can do lots of things with a little creativity. For example, you might pass the mean or median value of a series.

```
In [113]: 1 #Replace with mean
          2 df_missing.fillna(df_missing.mean())
```

```
Out[113]:
```

	name	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.000000	6	160	110.000000	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.733333	6	160	110.000000	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.800000	4	108	93.000000	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.400000	6	258	104.333333	3.08	3.215	19.44	1	0	3	1

```
In [ ]: 1
```